



Act-Aware Slot-Value Predicting in Multi-Domain Dialogue State Tracking

Ruolin Su, Ting-Wei Wu, Biing-Hwang Juang

Georgia Institute of Technology, USA

{ruolinsu, waynewu}@gatech.edu, juang@ece.gatech.edu

Abstract

As an essential component in task-oriented dialogue systems, dialogue state tracking (DST) aims to track human-machine interactions and generate state representations for managing the dialogue. Representations of dialogue states are dependent on the domain ontology and the user's goals. In several task-oriented dialogues with a limited scope of objectives, dialogue states can be represented as a set of slot-value pairs. As the capabilities of dialogue systems expand to support increasing naturalness in communication, incorporating dialogue act processing into dialogue model design becomes essential. The lack of such consideration limits the scalability of dialogue state tracking models for dialogues having specific objectives and ontology. To address this issue, we formulate and incorporate dialogue acts, and leverage recent advances in machine reading comprehension to predict both categorical and non-categorical types of slots for multi-domain dialogue state tracking. Experimental results show that our models can improve the overall accuracy of dialogue state tracking on the MultiWOZ 2.1 dataset, and demonstrate that incorporating dialogue acts can guide dialogue state design for future task-oriented dialogue systems.

Index Terms: dialogue state tracking, dialogue acts, task-oriented dialogue, reading comprehension

1. Introduction

With the rising demand of automatic human-machine interactions for accomplishing service tasks via a natural language dialogue, task-oriented dialogue systems have been developed and widely applied nowadays. Typically, a task-oriented dialogue system consists of four components: automatic speech recognition (ASR), natural language understanding (NLU), dialogue management (DM), and natural language generation (NLG). Dialogue state tracking (DST) is the core function of the DM module which tracks human-machine interactions and generates state representations for managing the conversational flow in a dialogue. Specifically, the DST system takes the results of a speech recognizer and a natural language understander, combined with dialogue context as input to predict distributions over a set of pre-defined variables [1]. In a task-oriented dialogue, dialogue states are the probability distributions of user's goals in belief space until the current user utterance.

Representations of dialogue states are dependent on the domain ontology and the involved user's goals. Table 1 shows an example of state representations within a task-oriented service dialogue between a user and a system for service across train and hotel domains. Among services via a natural language dialogue, some have well-defined objectives to achieve, e.g. reserving hotels or booking train tickets, in which typical dialogue state representations are a set of (slot, value) pairs, e.g., (destination, cambridge) and (day, wednesday) in a train ticket book-

Table 1: An example of cross-domain dialogue with dialogue state representations and system dialogue acts in MultiWOZ 2.1

<p>USER: Hi, I am looking for a train that is going to cambridge and arriving there by 20:45, is there anything like that?</p> <p><i>Dialogue States:</i></p> <p>TRAIN: destination=cambridge, arriveby=20:45</p>
<p>SYSTEM: Where will you be departing from?</p> <p>USER: I am departing from Birmingham New Street.</p> <p><i>Dialogue States:</i></p> <p>TRAIN: destination=cambridge, arriveby=20:45, departure=birmingham new street</p> <p><i>Dialogue Acts: Inform, Request</i></p>
<p>SYSTEM: Can you confirm your desired travel day?</p> <p>USER: I would like to leave on Wednesday.</p> <p><i>Dialogue States:</i></p> <p>TRAIN: destination=cambridge, arriveby=20:45, departure=birmingham new street, day=wednesday</p> <p><i>Dialogue Acts: Request</i></p>
<p>SYSTEM: I have booked your train tickets, and your reference number is a9nhso9y.</p> <p>USER: Thanks so much. I would also need a place to stay. I am looking for something with 4 stars and has free WiFi.</p> <p><i>Dialogue States:</i></p> <p>TRAIN: destination=cambridge, arriveby=20:45, departure=birmingham new street, day=wednesday</p> <p>HOTEL: stars=4, internet=yes, type=hotel</p> <p><i>Dialogue Acts: OfferBooked</i></p>

ing service. The objective of DST is therefore to accurately estimate the user's goals in previous dialogue and to represent them as such slot-value pairs. Such slot-value representations are widely used in quite a few task-oriented dialogues, such as ATIS [2], DSTC2 [3], MultiWOZ 2.0/2.1 [4, 5], etc.

In a human dialogue, speech acts are illocutionary actions contained in utterances that change dialogue states [6]. Similarly, those representing the illocutions of utterances in a human-machine dialogue are known as dialogue acts. Generally speaking, dialogue acts are defined in domain ontology and serve the functions of conducting particular tasks, having potentials to guide user utterances and enhance the performance of DST as auxiliary inputs. As an example in Table 1, the "Request" act by the system can result in a dialogue state transformation in the "train" domain.

Early works on DST regard speech acts as noisy observations of dialogue acts to update dialogue states. Assuming fixed domain ontology, a line of generative methods [7, 8] are proposed to represent dialogue states at each turn by modeling the joint probabilities in a belief space, which costs enormous manual efforts and limits the scalability to multi-domain dialogues. More recent works represent dialogue states as a set

of slot-value pairs [9, 10], where discriminative models have proved their capabilities in tracking dialogue states, by modeling DST as multi-task classification [11, 12, 13, 14] or question-answering [15, 16, 17, 18] problems. Recent discriminative DST models estimate user’s goals directly from the dialogue context, ignoring the NLU module and dialogue acts. However, using dialogue acts not only helps reason how dialogue states are predicted, but also improves compatibility of the DST model to existing pipeline dialogue systems and scalability to new domains.

To incorporate dialogue acts in discriminative DST models, we propose an act-aware dialogue state tracker (**ADST**) to predict slot-value pairs for tracking dialogue states with reasoning and high accuracy. We utilize system dialogue acts because they are easier to acquire and are closely relevant to dialogue system design. Furthermore, we exploit advances in reading comprehension (RC) to extend our DST models for task-oriented dialogues with more free-form domain ontology. Inspired by the RC-based work on DST [17, 18], we formulate the DST problem as predicting values by querying with two types of slots: categorical slots and non-categorical slots. For categorical slots, we implement multiple-choice RC [19] to choose from a pre-defined set of limited values, e.g. the “*hotel parking*” slot with a provided value set $\{Yes, No, Don't\ Care, None\}$. For non-categorical slots, slots are firstly determined to be either *Don't Care*, *None* or a span extracted from dialogue context, and then the system takes a span-based RC approach to predict values as probabilities of start and end positions in the dialogue context, e.g. “20:45” for the “*train arriveby*” slot. With the formulation of DST as reading comprehension tasks, our model can predict slot-values without pre-defined value sets by extracting values directly from the dialogue context. We also take advantage of pre-trained ELMo embeddings [20] to learn word syntax and semantics in dialogue context. In short, our contributions are as follows. (1) We leverage dialogue acts to attend on slots which bring about accuracy improvement on DST; (2) we propose models for categorical and non-categorical slots formulating DST as RC tasks with scalability; (3) we show dialogue acts are related to and have impact on DST via ablation study and attention weight visualization. Code will be available at: <https://github.com/youlandasu/ACT-AWARE-DST>.

2. Methods

2.1. Problem Formulation

We denote the tokenized user utterance as u_t^{usr} and the tokenized agent utterance as u_t^{sys} at dialogue turn t . The context C_t of the given dialogue at turn t is defined as the concatenation of the previous agent and user utterances, i.e. $C_t = \{u_1^{sys}, u_1^{usr}, \dots, u_t^{sys}, u_t^{usr}\}$. C_t is analogous to a passage in reading comprehension where the model predicts answers. The sequence of system dialogue acts until turn t is $A_t = \{a_1, a_2, \dots, a_t\}$, where $a_i = \{a_i^1, \dots, a_i^{l_i}\}$ represents the number of l_i dialogue acts in turn i . For task-oriented dialogues in each domain $d \in D$, the domain ontology defines a set of slots $S^d = S_c^d \cup S_n^d$, where S_c^d and S_n^d are the sets of categorical and non-categorical slots without overlapping.

For categorical slots S_c^d , we construct a passage C_t , a question $q_{d,s} = \{d, s\}$ and options $O_{d,s} = \{V_c^s, none, dont_care\}$ as those in multiple-choice reading comprehension. V_c^s is the set of possible values in each slot $s \in S^d$. Specifically, a question consists of a domain name and a slot name, and options are a list of all possible values with two special values: *none*

and *dont_care*. For non-categorical slots S_n^d , the passage and the question would be the same but options are $O_{d,s} = \{span, none, dont_care\}$, which substitutes with a span instead of a pre-defined set of values. If the option of “*span*” applies, predicting values from the dialogue context is equivalent to querying passage C_t with $q_{d,s}$ to find the best matched span in the passage.

Figure 1 shows the overall architecture of our proposed models, mainly consisting of a context encoder, one attention layer attending dialogue context, another attention layer attending system dialogue acts, and two similarity measure modules.

2.2. Encoding

Context Encoding. Denoting the i -th word in C_t as c_i , we combine word embedding, role embedding and binary exact match features for c_i as the input to the encoder. Specifically, word embeddings $W^e = [W^{ELMo}; W^{Char}]$ are formed by concatenating W^{ELMo} which are ELMo [20] pre-trained word embeddings and W^{Char} which are character-level tokens encoded by CNNs [21]. $W^e \in \mathbb{R}^{|C_t| \times w}$, where w is the sum of word embedding dimension and character embedding dimension while $|C_t|$ is the number of word tokens in the dialogue context. Role embeddings W^r are symbols “*SYS*” or “*USER*” to distinguish the system and the user in a dialogue. Exact match features W^{exact} are binary vectors that reflect where each pre-defined value in the domain ontology is shown in the previous dialogue context. The final input to the encoder is denoted as $W^{e'} = [W^e; W^r; W^{exact}]$. Then we use a bidirectional GRU [22] to encode $W^{e'}$ into $X^e \in \mathbb{R}^{|C_t| \times w}$, i.e. $X^e = BiGRU(W^{e'})$, and let the number of bidirectional hidden states be w , such that the dimension of encoder’s output is the same as that of W^e .

Dialogue Act and Slot-Value Encoding. The input dialogue act embeddings W^{act} are concatenated word embeddings of system dialogue acts in previous turns. We construct similar word embeddings $q^d, q^s \in \mathbb{R}^w$ for each domain and slot, respectively. For each dialogue, there are M domain-slot combinations corresponding to M questions. As for categorical slots, we sum q^d and q^s with each value embedding w^v in the value set V_c^s in addition to “*none*”, “*dont_care*”, and stack them together to construct option embeddings $P_{d,s}^e \in \mathbb{R}^{(N+2) \times w}$, i.e. $P_{d,s,i}^e = q^d + q^s + w_i^v$, where N is the number of values in the value set V_c^s , and $w_i^v \in \mathbb{R}^w$ is the i -th value in V_c^s .

2.3. Dialogue Context Attention

We implement an attention model similar to [21] that computes attention weights between dialogue context and slots. Assuming R, S are two matrices with the same columns h , the attention function is defined as:

$$Attention_k(R, S)_j = softmax_i([R_{i,:}; S_{j,:}; R_{i,:} \circ S_{j,:}] \cdot k) \quad (1)$$

where $k \in \mathbb{R}^{3h}$ is a trainable vector, \circ is element-wise multiplication, $[\cdot]$ is vector concatenation across column. According to the above definition, the attention weights is computed as $\alpha_{d,s}^{k_1} = Attention_{k_1}(X^e, q^d + q^s)$, and $\alpha_{d,s}^{k_1} \in \mathbb{R}^{|C_t|}$. Then the attended slot vector over dialogue context is $Q_{d,s}^c = (X^e)^T \cdot \alpha_{d,s}^{k_1}$. As a result, $Q_{d,s}^c \in \mathbb{R}^w$ is the output of slot embeddings which are dependent on the dialogue context.

2.4. Dialogue Act Attention

In order to fuse information from system dialogue acts, we compute an attended slot vector over dialogue acts following Equation 1. The attention weight of a querying slot attending to acts

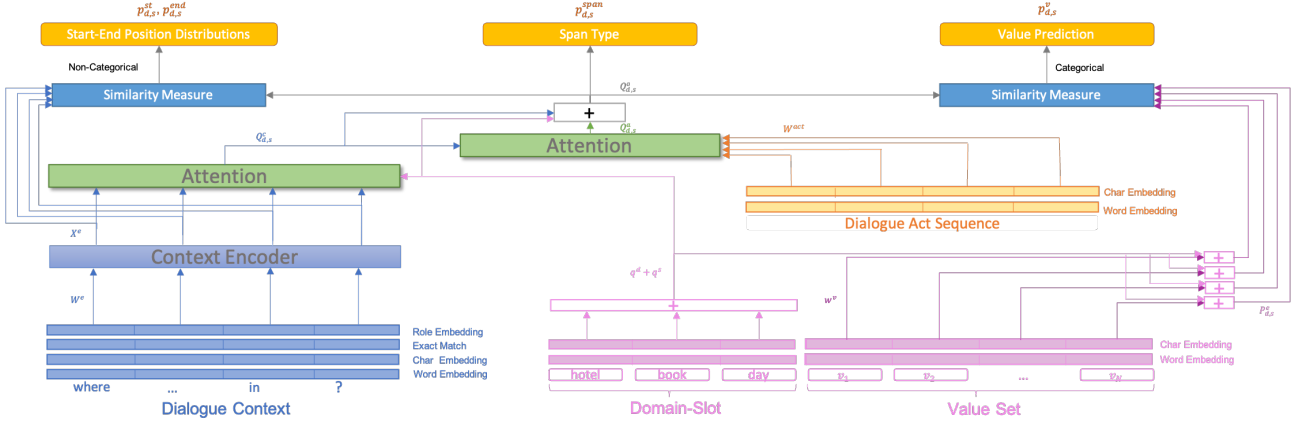


Figure 1: *Act-Aware Model Architecture for Dialogue State Tracking.* The domain-slot embeddings attend to both encoded dialogue context and previous system dialogue acts. For categorical slots, the similarity scores between possible values and an attended domain-slot is measured to choose from a set of values. For non-categorical slots, a span type is determined and the similarity scores between dialogue context and an attended domain-slot is measured to predict the position of a span.

is given as $\alpha_{d,s}^{k_2} = \text{Attention}_{k_2}(W^{act}, Q_{d,s}^c)_{\{d,s\}} \in \mathbb{R}^{|A_t|}$, where $W^{act} \in \mathbb{R}^{|A_t| \times w}$ is the word embedding of the dialogue acts and $|A_t|$ is the total number of system dialogue acts in previous turns. After that, we obtain a slot vector attended by system acts as $Q_{d,s}^a = (W^{act})^T \cdot \alpha_{d,s}^{k_2} \in \mathbb{R}^w$.

Then $Q_{d,s}^a$ is combined with $Q_{d,s}^c$ and the original slot embedding, i.e. $Q_{d,s}^o = Q_{d,s}^c + Q_{d,s}^a + q^d + q^s$, which can be regarded as the final slot embeddings dependent on both the previous dialogue acts and the context. Such that dialogue acts are incorporated in slots.

2.5. Value Classification for Categorical Slots

For each categorical slot, a value is to be selected from a pre-defined value set V_c^s . Inspired by [19], we compute probability of a value by calculating the bi-linear similarity between possible options $P_{d,s}^e$ and a final slot representations $Q_{d,s}^o$:

$$p_{d,s}^v = \text{softmax}(P_{d,s}^e \Theta^v Q_{d,s}^o) \quad (2)$$

where Θ^v is a trainable weight matrix. Denoting $y_{d,s}^v$ as true values for each categorical slots in the dialogue, then the cross entropy loss for the value prediction is calculated as:

$$L_v = \sum_t \sum_{d,s} \text{CrossEntropy}_t(p_{d,s}^v, y_{d,s}^v) \quad (3)$$

2.6. Span Prediction for Non-Categorical Slots

For non-categorical slots, we first decide the type of span from one of the following options: a span can be extracted from the dialogue context, “*dont_care*”, or “*none*”. The probability of the span type is calculated by $p_{d,s}^{span} = \text{softmax}(FFN_{type}(Q_{d,s}^o))$, where FFN_{type} represents a feed-forward neural network with output dimension of 3. Then we predict the probability distribution of start and end positions in the dialogue context with the following similarity functions:

$$p_{d,s}^{st} = \text{softmax}(FFN_{c_1}(X^e) \Theta^s Q_{d,s}^o) \quad (4)$$

$$p_{d,s}^{end} = \text{softmax}(FFN_{c_2}(X^e) \Theta^e Q_{d,s}^o) \quad (5)$$

where FFN_{c_1}, FFN_{c_2} are one-layer feed-forward networks with output dimensions of w , and Θ^s, Θ^e are two trainable

weight metrics for predicting the *start* and the *end* position, respectively. Denoting $y_{d,s}^{span}$ as the encoded true label type, the loss function for span type prediction is:

$$L_{type} = \sum_t \sum_{d,s} \text{CrossEntropy}_t(p_{d,s}^{span}, y_{d,s}^{span}) \quad (6)$$

Then we denote binary vectors of the true start and end positions as $y_{d,s}^{st}$ and $y_{d,s}^{end}$, respectively. The cross entropy loss for predicting span positions is as following:

$$L_s = \sum_t \sum_{d,s} \text{CrossEntropy}_t(p_{d,s}^{st}, y_{d,s}^{st}) + \sum_t \sum_{d,s} \text{CrossEntropy}_t(p_{d,s}^{end}, y_{d,s}^{end}) \quad (7)$$

Finally, the total loss is defined as $L = L_v + L_{type} + L_s$.

3. Experiments

3.1. Dataset

We train and evaluate our models on the MultiWOZ 2.1 dataset [5], which is a cross-domain task-oriented dialogue dataset collected from 7 domains containing over 10,000 multi-turn dialogues. MultiWOZ 2.1 contains 13 dialogue acts and 30 (*domain, slot*) combinations with hundreds of possible values. We split the dataset into training, development and test set the same as the original setting, and only use 5 most frequent domains in the dataset: {*restaurant, hotel, train, attraction, taxi*}.

3.2. Training Details

For the input context embeddings, we combine ELMo word embeddings with a length of 512, character embeddings with a length of 100, role embeddings with a length of 128 and one-hot exact matching features indicating occurrences of pre-defined values. For the context encoder, we use a one-layer bi-directional GRU with hidden units of the same length as ELMo embeddings combining character embeddings. ReLu [23] activation is used for all feed-forward layers. The learning rate is 0.001 with the ADAM optimizer [24] and the batch size is 24 in our joint training on 30 categorical and non-categorical slots.

Table 2: Joint and Slot Goal Accuracy on MultiWOZ 2.1

Model	Joint Goal Accuracy	Slot Goal Accuracy
w/o Non-Categorical Slots:		
DS-DST picklist [17]	53.30	-
DSTQA w/o span [18]	51.44	97.24
CHAN [25]	58.55	98.14
ADST (Ours) all categorical	56.70	97.71
w/ Non-Categorical Slots:		
STARC [15]	49.48	-
DS-DST [17]	51.21	-
DSTQA w/ span [18]	51.36	97.22
ADST (Ours) hybrid	56.12	97.62

3.3. Results

Table 2 lists the experimental results on MultiWOZ 2.1 test set, where the joint goal accuracy is the average accuracy of predicting all slot-values for a turn correctly, while the slot goal accuracy is the average accuracy of predicting the value of a slot correctly. We compare our models with: (1) DS-DST [17] which uses BERT-based [26] RC approaches to handle different slot types jointly, (2) DSTQA [18] which constructs slots with domain ontology for a RC-based model enhanced by a dynamic knowledge graph, (3) CHAN [25] which fine-tunes BERT in a hierarchical attention network to leverage relevant dialogue context, and 4) STARC [15] which pre-trains on RC dataset then fine-tunes on task dialogues to alleviate data scarcity problems.

We first train a model taking all slots as categorical, comparing it with the categorical-only models: DS-DST picklist, DSTQA w/o span, and CHAN. We achieve 56.70% joint goal accuracy and 97.71% slot goal accuracy on categorical-only slot-value predictions, which is close to the state-of-the-art joint and slot goal accuracy on the MultiWOZ 2.1 test set. In contrast to the current state-of-the-art model, our model is lightweight and can be scaled to non-categorical slots. In our hybrid model, we take all number- or time-related slots as non-categorical, whereas other slots as categorical, and train all slots jointly. We compare our results with those of STARC, DS-DST and DSTQA w/ span. We obtain outperformed results of 56.12% and 97.62% on joint and slot goal accuracy, due to the advantages of exploiting RC approaches and using system dialogue acts as auxiliary inputs.

Table 3: Ablation Study on MultiWOZ 2.1 Dev Set

Model	Dev Joint	Dev Slot
ADST (Ours) All Categorical	56.89	97.77
- w/o Dialogue Acts	53.95(-2.94)	97.53(-0.24)
ADST (Ours) All Non-Categorical	45.50	96.74
- w/o Dialogue Acts	44.48(-1.02)	96.65(-0.09)

Ablation. We evaluate our models on categorical and non-categorical slots either attending previous system dialogue acts to slots or not. Table 3 presents the ablation study results with regard to dialogue acts attention evaluated on the MultiWOZ 2.1 dev set. For our ADST model trained on all categorical slots, removing dialogue act attention layer drops 2.94% on the joint goal accuracy, and drops 0.24% on the slot goal accuracy. For our ADST model trained on all non-categorical slots, ignoring dialogue acts brings about 1.02% reduction of joint goal accuracy and 0.09% reduction of the slot goal accuracy. The ablation indicates that our models take advantage from input system di-

alogue acts for predicting slot-value pairs. We also observe that incorporating system dialogue acts into the representations of slots improves the performance of predicting values from the value set by 5% on joint slot accuracy, but has less impact on span predictions of non-categorical slots.




Figure 2: Visualization of attention weights on the dialogue act attention layer of our model for categorical slots.

Visualization of Attention Weights. To investigate how dialogue acts impact slot-value predictions, we feed a simulated system dialogue act sequence into our trained categorical-only model, and visualize the attention weights between dialogue acts and individual slots in Figure 2. We observe that the “Request” act attends the most weights to slots, whereas the “Welcome” act attends the least weights. That may be because a “Request” act guides a user to give more state-related information in a task-oriented dialogue, but a “Welcome” act is not targeted to any dialogue states. Dialogue acts like “NoOffer”, “Reqmore” yield higher attention weights, probably because they are more targeted to specific dialogue states than more general dialogue acts like “Select” or “Inform”. Note that the last act usually has comparably higher weights on all slots since it is most relevant to the current dialogue state. Results show that additional attention weights brought by including system dialogue acts suggest correlations between slots and dialogue acts, such that brings about performance improvement on DST.

4. Conclusion

We propose an act-aware method for multi-domain DST by incorporating system dialogue acts and dialogue context in previous turns to predict slot-value pairs up to the current turn. Our models combine dialogue acts, dialogue context and domain ontology, and leverages reading comprehension approaches to predict slots for both categorical and non-categorical slots. Experimental results show that attentions on both dialogue acts and dialogue context not only improve the joint goal accuracy on MultiWOZ 2.1, but also expand capacities of dialogue systems on reasoning how dialogue states are guided and transformed. The analysis and visualizations indicate that our model is able to use information of system dialogue acts to improve DST on specific slots. We believe that this idea of leveraging dialogue acts in discriminative DST models will improve their scalability for new domains and will contribute to the design of task-oriented dialogue systems for new services.

5. References

- [1] J. D. Williams and S. Young, “Partially observable markov decision processes for spoken dialog systems,” *Computer Speech & Language*, vol. 21, no. 2, pp. 393–422, 2007.
- [2] C. T. Hemphill, J. J. Godfrey, and G. R. Doddington, “The atis spoken language systems pilot corpus,” in *Speech and Natural Language: Proceedings of a Workshop Held at Hidden Valley, Pennsylvania, June 24-27, 1990*, 1990.
- [3] M. Henderson, B. Thomson, and J. D. Williams, “The second dialog state tracking challenge,” in *Proceedings of the 15th annual meeting of the special interest group on discourse and dialogue (SIGDIAL)*, 2014, pp. 263–272.
- [4] P. Budzianowski, T.-H. Wen, B.-H. Tseng, I. Casanueva, S. Ultes, O. Ramadan, and M. Gašić, “Multiwoz—a large-scale multi-domain wizard-of-oz dataset for task-oriented dialogue modelling,” *arXiv preprint arXiv:1810.00278*, 2018.
- [5] M. Eric, R. Goel, S. Paul, A. Kumar, A. Sethi, P. Ku, A. K. Goyal, S. Agarwal, S. Gao, and D. Hakkani-Tur, “Multiwoz 2.1: A consolidated multi-domain dialogue dataset with state corrections and state tracking baselines,” *arXiv preprint arXiv:1907.01669*, 2019.
- [6] J. R. Searle and D. Vanderveken, “Speech acts and illocutionary logic,” in *Logic, thought and action*. Springer, 1985, pp. 109–132.
- [7] D. Bohus and A. Rudnicky, “A “k hypotheses+ other” belief updating model,” 2006.
- [8] S. Young, M. Gašić, B. Thomson, and J. D. Williams, “Pomdp-based statistical spoken dialog systems: A review,” *Proceedings of the IEEE*, vol. 101, no. 5, pp. 1160–1179, 2013.
- [9] N. Mrkšić, D. O. Séaghdha, T.-H. Wen, B. Thomson, and S. Young, “Neural belief tracker: Data-driven dialogue state tracking,” *arXiv preprint arXiv:1606.03777*, 2016.
- [10] A. Rastogi, X. Zang, S. Sunkara, R. Gupta, and P. Khaitan, “Towards scalable multi-domain conversational agents: The schema-guided dialogue dataset,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 05, 2020, pp. 8689–8696.
- [11] C.-S. Wu, A. Madotto, E. Hosseini-Asl, C. Xiong, R. Socher, and P. Fung, “Transferable multi-domain state generator for task-oriented dialogue systems,” *arXiv preprint arXiv:1905.08743*, 2019.
- [12] V. Zhong, C. Xiong, and R. Socher, “Global-locally self-attentive dialogue state tracker,” *arXiv preprint arXiv:1805.09655*, 2018.
- [13] A. Rastogi, D. Hakkani-Tür, and L. Heck, “Scalable multi-domain dialogue state tracking,” in *2017 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*. IEEE, 2017, pp. 561–568.
- [14] H. Shi, T. Ushio, M. Endo, K. Yamagami, and N. Horii, “A multichannel convolutional neural network for cross-language dialog state tracking,” in *2016 IEEE Spoken Language Technology Workshop (SLT)*. IEEE, 2016, pp. 559–564.
- [15] S. Gao, S. Agarwal, T. Chung, D. Jin, and D. Hakkani-Tur, “From machine reading comprehension to dialogue state tracking: Bridging the gap,” *arXiv preprint arXiv:2004.05827*, 2020.
- [16] S. Gao, A. Sethi, S. Agarwal, T. Chung, and D. Hakkani-Tur, “Dialog state tracking: A neural reading comprehension approach,” *arXiv preprint arXiv:1908.01946*, 2019.
- [17] J.-G. Zhang, K. Hashimoto, C.-S. Wu, Y. Wan, P. S. Yu, R. Socher, and C. Xiong, “Find or classify? dual strategy for slot-value predictions on multi-domain dialog state tracking,” *arXiv preprint arXiv:1910.03544*, 2019.
- [18] L. Zhou and K. Small, “Multi-domain dialogue state tracking as dynamic knowledge graph enhanced question answering,” *arXiv preprint arXiv:1911.06192*, 2019.
- [19] G. Lai, Q. Xie, H. Liu, Y. Yang, and E. Hovy, “Race: Large-scale reading comprehension dataset from examinations,” *arXiv preprint arXiv:1704.04683*, 2017.
- [20] M. E. Peters, M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, and L. Zettlemoyer, “Deep contextualized word representations,” *arXiv preprint arXiv:1802.05365*, 2018.
- [21] M. Seo, A. Kembhavi, A. Farhadi, and H. Hajishirzi, “Bidirectional attention flow for machine comprehension,” *arXiv preprint arXiv:1611.01603*, 2016.
- [22] K. Cho, B. Van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, “Learning phrase representations using rnn encoder-decoder for statistical machine translation,” *arXiv preprint arXiv:1406.1078*, 2014.
- [23] V. Nair and G. E. Hinton, “Rectified linear units improve restricted boltzmann machines,” in *Icml*, 2010.
- [24] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” *arXiv preprint arXiv:1412.6980*, 2014.
- [25] Y. Shan, Z. Li, J. Zhang, F. Meng, Y. Feng, C. Niu, and J. Zhou, “A contextual hierarchical attention network with adaptive objective for dialogue state tracking,” *arXiv preprint arXiv:2006.01554*, 2020.
- [26] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “Bert: Pre-training of deep bidirectional transformers for language understanding,” *arXiv preprint arXiv:1810.04805*, 2018.