



Sequential End-to-End Intent and Slot Label Classification and Localization

Yiran Cao¹, Nihal Potdar¹, Anderson R. Avila²

¹University of Waterloo, Canada

²Huawei Noah's Ark Lab, Canada

{yiran.cao, nihal.potdar}@waterloo.edu, anderson.avila@huawei.com

Abstract

Human-computer interaction (HCI) is significantly impacted by delayed responses from a spoken dialogue system. Hence, end-to-end (e2e) spoken language understanding (SLU) solutions have recently been proposed to decrease latency. Such approaches allow for the extraction of semantic information directly from the speech signal, thus bypassing the need for a transcript from an automatic speech recognition (ASR) system. In this paper, we propose a compact e2e SLU architecture for streaming scenarios, where chunks of the speech signal are processed continuously to predict intent and slot values. Our model is based on a 3D convolutional neural network (3D-CNN) and a unidirectional long short-term memory (LSTM). We compare the performance of two alignment-free losses: the connectionist temporal classification (CTC) method and its adapted version, namely connectionist temporal localization (CTL). The latter performs not only the classification but also localization of sequential audio events. The proposed solution is evaluated on the Fluent Speech Command dataset and results show our model ability to process incoming speech signal, reaching accuracy as high as 98.97 % for CTC and 98.78 % for CTL on single-label classification, and as high as 95.69 % for CTC and 95.28 % for CTL on two-label prediction.

Index Terms: spoken language understanding, human-computer interaction, low-latency

1. Introduction

Spoken language understanding (SLU) aims at extracting structured semantic representations, such as intent and slots, from the speech signal [1]. These representations are crucial to enable speech as the primary mode of human-computer interaction (HCI) [2]. Traditional SLU solutions rely on the text transcription generated by an automatic speech recognition (ASR) module, followed by a natural language understanding (NLU) system, responsible for extracting semantics from the ASR output [3]. As described in [4], in such scenarios the ASR typically operates on chunks of the incoming speech signal and outputs the transcript for each segment. The NLU then waits until all the speech segments are transcribed before processing the ASR output. This has significant latency implications. Another issue is that each module is trained and optimized separately. While the ASR optimization aims at minimizing word error rate, the NLU is often optimized on clean text with the assumption of error-less transcriptions from the ASR [5]. Besides, this approach provides a cumulative error that propagates from each module, adding up to the overall SLU error.

Recently, we have witnessed an increasing interest in reducing the latency of the SLU task. Low-latency leads to more naturalness while interacting with a computer system and can ultimately improve the user experience (UX) [6]. To this end, a handful of studies have specifically addressed the problem. In

[4], the authors proposed a recurrent neural network (RNN) for processing the output of an ASR system in an online fashion. Their streaming SLU solution is based on an online NLU that processes word sequences of arbitrary length and incrementally provides multiple intent predictions. Similarly, the authors in [7] propose an RNN-based model that jointly performs online intent detection and slot filling as input word embeddings arrive. Results show that the joint training model provides high accuracy for intent detection and language modeling with a small degradation on slot filling compared to the independent training models. Although these approaches show reasonable performance, they rely on the strong assumption of error-less transcriptions from the ASR as their NLU system is often trained on clean text. Moreover, they can not be considered end-to-end (e2e) solutions as their models are based on the ASR transcription.

To mitigate this, other studies have proposed the extraction of semantic information directly from audio. For example, several e2e SLU encoder-decoder architectures are investigated in [8]. The authors showed that better performance is achieved when an e2e SLU solution that performs domain, intent, and argument prediction is jointly trained with an e2e ASR model that learns to generate transcripts from the same input speech. Another recent study introduces the Fluent Speech Command (FSC) dataset [9]. The authors present a pre-training strategy for e2e SLU models. Their approach is based on using ASR targets, such as words and phonemes, that are used to pre-train the initial layers of their final model. These classifiers once trained are discarded and the embeddings from the pre-trained layers are used as features for the SLU task. Improved performance on large and small SLU training sets was achieved with the proposed pre-training approach. Similarly, in [10], the authors also proposed to fine-tune the lower layers of an end-to-end CNN-RNN based model that learns to predict graphemes. This pre-trained acoustic model is optimized with the connectionist temporal classification (CTC) loss and then combined with a semantic model to predict intents.

The aforementioned research efforts have been either on developing online NLU or non-streamable e2e SLU. In the light of that, investigating a complete end-to-end low-latency streaming SLU solution is necessary. In this paper, we propose a compact e2e streamable SLU solution that (1) eliminates the need for an ASR module with (2) an online architecture that provides intent and slot predictions while processing incoming speech signals. To achieve that, a 3-dimensional convolutional neural network (3D-CNN) combined with a unidirectional long-short term memory (LSTM) is explored. We compare two alignment-free loss functions: the CTC method and its adaptation, namely the connectionist temporal localization (CTL) function. Both methods will be discussed in section 2. We use the FSC dataset to perform our experiments and results show our model achieving accuracy as high as 98.97 % for single intent+slot classification.

x_1	x_2	x_3	x_4	x_5	x_6	x_{17}	x_8	x_9	x_{10}	x_{11}	x_{12}	x_{13}	x_{14}	x_{15}	x_{16}	x_{17}	x_{18}	x_{19}	x_{20}
t	u	r	r		n		o	o	n		t	h	e			c	c	a	r
turn							on				the					car			
Activate																car			

Figure 1: CTC alignment for $X = [x_1, x_2, \dots, x_n]$, output $Y_{asr} = [t, u, r, n, \rightarrow, o, n, \rightarrow, t, h, e, \rightarrow, c, a, r]$ and $Y_{slu} = [activate, car]$.

cation and 95.69 % for multiple intent+slot detection.

The remainder of this paper is organized as follows. In Section 2, we discuss sequential labeling. Section 3 presents the proposed architecture. Section 4 describes our experimental setup and Section 5 presents our experimental evaluation. Section 6 gives the conclusion.

2. Sequential Labeling

2.1. Streaming Spoken Language Understanding

In a streaming e2e SLU scenario, given the input $X = [x_1, x_2, \dots, x_n]$ of acoustic features of length N and the corresponding sequence of semantic outputs $Y = [y_1, y_2, \dots, y_u]$ of length U , the precise alignment of X and Y is not known and typically $N > U$. Unlikely ASR, for e2e SLU the gap between input and output length is higher as the semantic label prediction is conditioned to a larger input context. The period of silence in the audio tends to even increase this gap. The goal is to learn the distribution $P(Y|X, \theta)$. However, different from the non-streaming scenario, predictions are made for a given timestep, t , with the model incrementally predicting multiple output intents (or blank symbols for silence), while accounting for the context dependency from previous predictions.

2.2. Connectionist Temporal Classification

The CTC method was first motivated to train RNNs of unsegmented data [11]. Previous to CTC, training RNNs required prior segmentation of the input sequence. For that, each input segment was labeled and the RNN was trained to predict an output for each segment at a time. This required a post-processing step to consolidate the output predictions into a sequence of labels [11]. With CTC, however, prior segmentation is no longer needed as the method allows a sequence-to-sequence mapping free of alignment. For acoustic modeling, for example, CTC automatically learns the alignments between the input sequence of acoustic frames and the respective sequence of output labels [12]. As illustrated in Figure 1, CTC defines the so-called naive alignment by matching the input and output length adding the blank tokens (), and repeating output predictions. After this, blank tokens are removed and repeated predictions are collapsed.

Formally, the conditional probability of a single alignment (or path), α , is the product of the probabilities of observing α_t at time t and can be represented as

$$P(\alpha|X, \theta) = \prod_{t=1}^T P(\alpha_t|X, \theta) \quad (1)$$

where α_t represents a given label. Because $P(\alpha|X)$ defines mutually exclusive paths, the conditional probability for a sequence output is given by the sum of the probabilities of all paths corresponding to it:

$$P(Y|X, \theta) = \sum_{\alpha \in A_{X,Y}} \prod_{t=1}^T P(\alpha_t|X, \theta) \quad (2)$$




Figure 2: Streaming intent + slot approach versus a non-streaming one. Note that slot has a higher prediction rate than intent (3 times faster in our experiments) as it depends on less temporal speech context.

where $A_{X,Y}$ is the set of all valid alignments. The CTC loss is then defined as

$$L_{CTC}(X, Y) = -\log \sum_{\alpha \in A_{X,Y}} \prod_{t=1}^T P(\alpha_t|X, \theta) \quad (3)$$

CTC considers no dependency between previous time steps which allows for frame-wise gradient propagation, but limits the possibility of learning sequential dependencies [13].

2.3. Connectionist Temporal Localization

CTL was first introduced in [14] for classification and localization of sound event occurrences in an audio stream. CTC has been previously applied for sound event detection (SED), however, it was found that the method presented the so-called "peak clustering," especially for long events [15]. In such cases, onset and offset labels will result in a peak of the frame-probabilities. During training, because the adjacent onset and offset labels of long events occur next to each other, CTC may interpret them as the existence of boundaries instead of the existence of an event. Because the CTC loss function focus on predicting sequential labels in the correct order regardless of any temporal constraints, the recurrent layer will derive onset and offset labels next to each other as it minimizes memory effort [14]. Moreover, the model detects event boundaries which leads to high frame-probabilities surrounding onset and offset events and remains inactive for the period that the event is on, even when changes in the acoustic features are observed.

Three improvements are proposed to overcome the peak clustering issue [14]. First, boundary probabilities are attained from network event probabilities using a "rectified delta" operator. This assures that the network predicts frame-wise probabilities of events and not of the event boundaries, which leads to different predictions for different acoustic features [14]. The event boundaries are calculated then as follows,

$$\begin{aligned} z_t(\dot{E}) &= \max[0, y_t(E) - y_{t-1}(E)] \\ z_t(\ddot{E}) &= \max[0, y_{t-1}(E) - y_t(E)] \end{aligned} \quad (4)$$

where $y_t(E)$ is the probability of the event E at frame t , whereas $z_t(\dot{E})$ and $z_t(\ddot{E})$ represent, respectively, onset and offset labels of event E at frame t .

Second, the boundary probabilities at each frame are considered mutually independent, which allows the overlap of sound events. The independence assumption eliminates the

Table 2: Results on multi-intent classification. The training set consists of 1 label and the testing set contains utterances with 1 and 2 labels, referred to as FSC-M1-Tst and FSC-M2-Tst.

Model	FSC-M1-Tst			FSC-M2-Tst		
	Intent	Slot	Intent+Slot	Intent	Slot	Intent+Slot
CTC	27.70	54.41	14.50	29.92	24.83	5.30
+ Joint CE	98.25	99.34	97.73	53.00	60.55	41.31
+ Pretrained ASR	99.15	99.76	98.97	56.04	64.21	44.30
CTL	12.76	55.54	3.60	0.14	25.05	0.05
+ Joint CE	89.05	95.35	85.68	48.37	80.14	41.28
+ Pretrained ASR	97.25	99.12	96.41	44.50	81.22	39.29
CTL + MIL	98.54	98.39	96.99	78.21	86.00	69.10
+ Pretrained ASR	99.18	99.57	98.78	66.17	96.77	64.89

4.2. Features

In this work, audio signals are sampled at 16 kHz. As acoustic features, 80-dimensional log Mel-Filterbank features are adopted [22]. To extract the Mel features, the audio signal is processed in frames of 320 samples (i.e., 20-ms window length), with a step size of 160 samples (that is, 10-ms hop-size). Global Cepstral Mean and Variance Normalization (CMVN) is applied, as a process that is commonly used with the aim to increase the robustness of ASR systems while mitigating the mismatch between training and testing data [23][24]

4.3. Experimental Settings

Our network was trained on mini-batches of 64 samples over a total of 200 epochs. In our experiments, we adopted learning rate of 0.0001 and dropout of 0.1. The optimizer used was Adam with weight decay of 0.2. We explored three strategies to optimize our model for the streaming scenario, including (1) using pure CTC or pure CTL as loss function; (2) using CTC or CTL jointly with CE. In the case of CTL, we also attempted to combine it with the multiple instance learning (MIL) technique [14]; and finally (3) adding a pretrained ASR to our model, optimized with the first layer with the CTC loss on character prediction. Note that CE and the MIL were only used during training. For the CE, we applied softmax to the last timestep and average the CTC and CE losses with a fixed weight of 0.6 and 0.4, respectively. Because the MIL is based on frame-wise probabilities just as CTL, their combination is straight-forward and consists of a simple weighted average of the two losses. The main difference is that the MIL aggregates the frame-probabilities into recording-level probabilities with a linear softmax pooling function [14]. Note that to ensure the streaming ability of our model, at the testing time only CTC or CTL are used. The third strategy consisted of pretraining an ASR model with the CTC loss. This procedure aimed at leveraging pre-trained embeddings by learning better representations from a large corpus. The model was trained to predict characters and only the first layer of our neural network was used. After training the ASR for 150 epochs, its weights were frozen and used with the entire recurrent neural network. In our experiments, we defined intent as the combination of action and object, which led to a total of 15 different intent labels. Location was defined as slot, which led to a total of 8 different slot labels. Our model was evaluated on 3 tasks: intent prediction; slot prediction; and intent+slot prediction. While intent predictions are attained using the outputs of the third layer, slot predictions rely on the output of the second layer, as illustrated in Figure 3.

Table 3: Results on multi-intent classification. The training set consists of 2 labels and the testing set contains utterances with 1 and 2 labels, referred to as FSC-M1-Tst and FSC-M2-Tst.

Model	FSC-M1-Tst			FSC-M2-Tst		
	Intent	Slot	Intent+Slot	Intent	Slot	Intent+Slot
CTC	22.83	0.00	0.00	22.43	25.03	4.24
+ Joint CE	97.49	99.23	96.88	95.16	97.77	93.73
+ Pretrained ASR	98.54	99.39	98.15	96.45	98.47	95.69
CTL	11.83	55.54	2.66	1.17	26.58	0.20
+ Joint CE	86.79	94.93	83.15	76.45	88.43	69.31
+ Pretrained ASR	95.38	99.10	94.67	87.58	97.33	85.50
CTL + MIL	96.70	98.89	95.75	93.14	96.83	90.30
+ Pretrained ASR	98.65	99.68	98.36	95.78	99.50	95.28

5. Experimental Evaluation

The performance of the proposed architecture is investigated in three different experiments. In the first one, we evaluate our model in a non-streaming scenario. The model is based on the pre-trained ASR and the architecture presented in Section 3. Only CE loss is used for optimizing our model. This experiment aims to compare our network with non-streamable e2e SLU solutions proposed in the literature. Results are reported in Table 1 and it shows our model achieving high accuracy in the three tasks and outperforming the other three baselines.

In the second experiment, we tackle a more challenging setting and our model is evaluated under the streaming regime. As described in Table 2, the model is trained with a single label (i.e. intent, slot, or intent+slot) and tested on one (FSC-M1-Tst) and two labels (FSC-M2-Tst). Results show that using the CE loss during training is a key strategy to boost the performance of both CTC and CTL. The use of a pretrained ASR is also crucial for achieving high accuracy. Also, combining CTL with MIL yields better accuracy overall. Identifying intent and slot correctly at the same time is difficult and provides lower accuracy when compared to the single task (i.e. single intent or single slot). Detecting 2 intents have a detrimental impact on the performance of our model. To overcome this, in the last experiment, our model is trained on two labels and tested on one and two labels as well. Results are presented on Table 3. We can observe that training on two labels benefits the performance of our model. Although CTC and CTL present comparable results, one has to consider that CTL has the potential for performing localization, whereas CTC is limited to predicting a sequence. Moreover, CTL allows for overlap events.

6. Conclusion

In this paper, we evaluated a compact spoken language understanding (SLU) model optimized with two alignment-free losses: the connectionist temporal classification (CTC) and the connectionist temporal localization (CTL). These losses allow the streaming capability on SLU, enabling the prediction of semantics based on incoming speech. In our first experiments, we showed that our model can achieve accuracy as high as 99.26 % on non-streaming settings while predicting intent and slot. For streaming scenarios, the proposed model can achieve accuracy of 98.97 % for CTC and 98.78 % for CTL on single label prediction. We also showed that our model is able to perform sequential labeling without compromising performance when it is presented with multiple utterances and targets during training and accuracy is as high as 95.69 % for CTC and 95.28 % for CTL on two label prediction. As future work, we plan to investigate the capability of our model to precisely locate and decode the semantic span within an utterance.

7. References

- [1] M. Mhiri, S. Myer, and V. S. Tomar, “A low latency asr-free end to end spoken language understanding system,” *arXiv preprint arXiv:2011.04884*, 2020.
- [2] Y. He, T. N. Sainath, R. Prabhavalkar, I. McGraw, R. Alvarez, D. Zhao, D. Rybach, A. Kannan, Y. Wu, R. Pang *et al.*, “Streaming end-to-end speech recognition for mobile devices,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2019, pp. 6381–6385.
- [3] Y. Qian, R. Ubale, P. Lange, K. Evanini, V. Ramanarayanan, and F. K. Soong, “Spoken language understanding of human-machine conversations for language learning applications,” *Journal of Signal Processing Systems*, pp. 1–13, 2019.
- [4] P. G. Shivakumar, N. Kumar, P. Georgiou, and S. Narayanan, “Incremental online spoken language understanding,” *arXiv preprint arXiv:1910.10287*, 2019.
- [5] D. Serdyuk, Y. Wang, C. Fuegen, A. Kumar, B. Liu, and Y. Bengio, “Towards end-to-end spoken language understanding,” in *2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2018, pp. 5754–5758.
- [6] M. Porcheron, J. E. Fischer, S. Reeves, and S. Sharples, “Voice interfaces in everyday life,” in *proceedings of the 2018 CHI conference on human factors in computing systems*, 2018, pp. 1–12.
- [7] B. Liu and I. Lane, “Joint online spoken language understanding and language modeling with recurrent neural networks,” *arXiv preprint arXiv:1609.01462*, 2016.
- [8] P. Haghani, A. Narayanan, M. Bacchiani, G. Chuang, N. Gaur, P. Moreno, R. Prabhavalkar, Z. Qu, and A. Waters, “From audio to semantics: Approaches to end-to-end spoken language understanding,” in *2018 IEEE Spoken Language Technology Workshop (SLT)*. IEEE, 2018, pp. 720–726.
- [9] L. Lugosch, M. Ravanelli, P. Ignoto, V. S. Tomar, and Y. Bengio, “Speech model pre-training for end-to-end spoken language understanding,” *arXiv preprint arXiv:1904.03670*, 2019.
- [10] Y.-P. Chen, R. Price, and S. Bangalore, “Spoken language understanding without speech recognition,” in *2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2018, pp. 6189–6193.
- [11] A. Graves, S. Fernández, F. Gomez, and J. Schmidhuber, “Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks,” in *Proceedings of the 23rd international conference on Machine learning*, 2006, pp. 369–376.
- [12] Y. Miao, M. Gowayyed, X. Na, T. Ko, F. Metze, and A. Waibel, “An empirical exploration of ctc acoustic models,” in *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2016, pp. 2623–2627.
- [13] U. Kamath, J. Liu, and J. Whitaker, *Deep learning for nlp and speech recognition*. Springer, 2019, vol. 84.
- [14] Y. Wang and F. Metze, “Connectionist temporal localization for sound event detection with sequential labeling,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2019, pp. 745–749.
- [15] —, “A first attempt at polyphonic sound event detection using connectionist temporal classification,” in *2017 IEEE international conference on acoustics, speech and signal processing (icassp)*. IEEE, 2017, pp. 2986–2990.
- [16] J. Kim, K. P. Truong, G. Englebiene, and V. Evers, “Learning spectro-temporal features with 3d cnns for speech emotion recognition,” in *2017 Seventh International Conference on Affective Computing and Intelligent Interaction (ACII)*. IEEE, 2017, pp. 383–388.
- [17] S. Ganapathy and V. Peddinti, “3-d cnn models for far-field multi-channel speech recognition,” in *2018 IEEE international conference on acoustics, speech and signal processing (ICASSP)*. IEEE, 2018, pp. 5499–5503.
- [18] H. Meng, T. Yan, F. Yuan, and H. Wei, “Speech emotion recognition from 3d log-mel spectrograms with deep learning network,” *IEEE access*, vol. 7, pp. 125 868–125 881, 2019.
- [19] K. Rao, H. Sak, and R. Prabhavalkar, “Exploring architectures, data and units for streaming end-to-end speech recognition with rnn-transducer,” in *2017 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*. IEEE, 2017, pp. 193–199.
- [20] V. Panayotov, G. Chen, D. Povey, and S. Khudanpur, “Librispeech: an asr corpus based on public domain audio books,” in *2015 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2015, pp. 5206–5210.
- [21] Y. Tian and P. J. Gorinski, “Improving end-to-end speech-to-intent classification with reptile,” *arXiv preprint arXiv:2008.01994*, 2020.
- [22] S. Watanabe, T. Hori, S. Karita, T. Hayashi, J. Nishitoba, Y. Unno, N. E. Y. Soplin, J. Heymann, M. Wiesner, N. Chen *et al.*, “Espnet: End-to-end speech processing toolkit,” *arXiv preprint arXiv:1804.00015*, 2018.
- [23] P. Pujol, D. Macho, and C. Nadeu, “On real-time mean-and-variance normalization of speech recognition features,” in *2006 IEEE international conference on acoustics speech and signal processing proceedings*, vol. 1. IEEE, 2006, pp. I–I.
- [24] M. Ravanelli, T. Parcollet, and Y. Bengio, “The pytorch-kaldi speech recognition toolkit,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2019, pp. 6465–6469.